

Forecast Service - Architecture Document


1. Overview

The Forecast Service is a microservice designed for managing power plant production forecasts in an energy trading platform. It provides RESTful APIs for creating, updating, and retrieving forecast data, as well as calculating company-wide position aggregations across multiple power plants.


2. System Architecture

2.1 Architectural Pattern


The service follows **Clean Architecture** principles with a clear separation of concerns across four layers:



2.2 System Flow Diagram



2.3 Domain Model



datetime	CreatedAt	
datetime	UpdatedAt	

3. API Endpoints

3.1 Forecast Management

POST /api/forecasts

Create or update a forecast for a power plant.

Request Body:

```
{
  "powerPlantId": "guid",
  "forecastDateTime": "2026-01-16T12:00:00Z",
  "productionMWh": 150.5
}
```

Response (201 Created / 200 OK):

```
{
  "success": true,
  "data": {
    "id": "guid",
    "powerPlantId": "guid",
    "powerPlantName": "Turkey Power Plant",
    "country": "Turkey",
    "forecastDateTime": "2026-01-16T12:00:00Z",
    "productionMWh": 150.5,
    "createdAt": "2026-01-16T10:00:00Z",
    "updatedAt": "2026-01-16T10:00:00Z"
  },
  "error": null
}
```

Error Response (400 Bad Request / 404 Not Found):

```
{
  "success": false,
  "data": null,
```

```

"error": {
    "title": "Not Found",
    "detail": "Power plant with ID ... not found",
    "status": 404,
    "instance": "/api/forecasts",
    "timestamp": "2026-01-18T10:30:00Z"
}
}

```

GET /api/forecasts/{id}

Retrieve a specific forecast by ID.

GET /api/forecasts/power-plant/{powerPlantId}

Get all forecasts for a power plant within a date range.

Query Parameters:

- `startDate`: DateTime
- `endDate`: DateTime

3.2 Company Position

GET /api/companyposition/{companyId}

Calculate and retrieve the aggregated position for a company.

Query Parameters:

- `startDate`: DateTime
- `endDate`: DateTime

Response:

```
{
  "success": true,
  "data": {
    "companyId": "guid",
    "companyName": "Energy Trading Corp",
    "startDate": "2026-01-16T00:00:00Z",
    "endDate": "2026-01-17T00:00:00Z",
    "totalPositionMWh": 1250.75,
    "powerPlantPositions": [
      {
        "powerPlantId": "guid",
        "powerPlantName": "Turkey Power Plant",
        "country": "Turkey",
        "positions": [
          {
            "start": "2026-01-16T00:00:00Z",
            "end": "2026-01-17T00:00:00Z",
            "mwh": 1250.75
          }
        ]
      }
    ]
  }
}
```


```

        "totalProductionMWh": 450.25,
        "forecastCount": 24
    }
]
},
"error": null
}


```

4. Data Flow

4.1 Create/Update Forecast Flow



4.2 Get Company Position Flow



5. Event-Driven Architecture

5.1 PositionChanged Event

When a forecast is created or updated, the service emits a `PositionChangedEvent`:

```
{
  "CompanyId": "guid",
  "StartDate": "2026-01-16T00:00:00Z",
  "EndDate": "2026-01-17T00:00:00Z",
  "TotalPositionMWh": 1250.75,
  "EventTimestamp": "2026-01-16T10:30:00Z",
  "Reason": "Forecast Created"
}
```

Current Implementation: RabbitMQ message broker running in Docker container

- **Exchange:** `forecast.events` (topic exchange)
- **Routing Key:** `position.changed`
- **Message Format:** JSON serialized events
- **Management UI:** <http://localhost:15672> (admin/admin)
- **Connection:** Automatic reconnection with retry logic

Alternative Options: Azure Service Bus, Apache Kafka, or AWS EventBridge

6. Data Persistence

6.1 Database Strategy

- **Database:** PostgreSQL 16
- **ORM:** Entity Framework Core 10.0
- **Migration Strategy:** Code-First with automatic migrations on startup
- **Connection Pooling:** Managed by Npgsql provider
- **Indexing Strategy:**
 - Composite index on (`PowerPlantId`, `ForecastDateTime`, `IsActive`) for fast forecast queries
 - Composite index on (`CompanyId`, `Country`) for power plant queries

6.2 Data Seeding

Initial data includes:

- 1 Company: "Energy Trading Corp"
- 3 Power Plants: Turkey, Bulgaria, Spain

7. Thread Safety & Concurrency

7.1 Thread Safety Considerations

1. **DbContext Scoping:** Each HTTP request gets its own scoped DbContext instance
2. **Repository Pattern:** Prevents shared state issues
3. **Async/Await:** All I/O operations are async for better scalability
4. **Optimistic Concurrency:** Can be added using row versioning if needed

7.2 Scalability

- **Stateless Design:** Service can be horizontally scaled
- **Database Connection Pooling:** Efficient resource utilization
- **Docker Support:** Easy container orchestration with Kubernetes or Docker Swarm

8. Security Considerations


8.1 Current Implementation

- **Result Pattern:** Type-safe error handling without exceptions
- **Model State Validation:** Nullable properties with `[Required]` and `[Range]` attributes
- **Global Exception Handler:** Middleware for consistent error responses
- **Custom Validation Error Factory:** Returns unified error format for validation failures
- **Structured error responses:** Consistent error format with codes and messages
- **CORS configuration:** Configured for cross-origin requests
- **Docker isolation:** Services isolated in separate containers

8.2 Production Recommendations

- Add JWT authentication
- Implement rate limiting
- Add API versioning
- Enable HTTPS only
- Add request logging and monitoring
- Implement API key authentication

9. Deployment Architecture



10. Monitoring & Observability

10.1 Recommended Metrics

- Request latency (p50, p95, p99)
- Error rates by endpoint
- Database query performance
- Event publishing success rate
- Active connections

10.2 Logging

- **Structured logging using Serilog:** Already implemented
- **Console Sink:** Real-time log output during development
- **File Sink:** Daily rolling logs stored in `logs/` directory with 30-day retention
- **Log levels:** Debug, Information, Warning, Error, Critical
- **Log format:** JSON structured format with timestamps, log levels, and context
- **Future Enhancement:** Correlation IDs for distributed request tracing

10.2 Modern C# Features

C# 14 field Keyword (Limited Usage):

- **BaseEntity only:** Used for UTC timestamp normalization in CreatedAt and UpdatedAt properties
- **Benefit:** Cleaner code by eliminating explicit backing fields for date normalization logic
- **Example:** `set => field = value.Kind == DateTimeKind.Unspecified ? DateTime.SpecifyKind(value, DateTimeKind.Utc) : value.ToUniversalTime();`

Validation Strategy - Data Annotations (ASP.NET Core Standard):

- **Approach:** Nullable properties with `[Required]` and `[Range]` data annotations
- **Rationale:**
 - Data annotations are the ASP.NET Core idiomatic way for input validation
 - Automatic validation before controller action execution
 - Consistent error responses through custom `InvalidModelStateResponseFactory`
 - Separates validation concerns from business logic
 - No exceptions during model binding - validation handled declaratively

Rationale for Pragmatic Approach:

- This is an **I/O-bound microservice** where network/database latency dominates performance
- Event publishing happens **occasionally** (on forecast changes), not in tight loops
- Advanced C# 14 features like `Span<T>` would add complexity without measurable benefits
- **Principle:** Favor readability, maintainability, and ASP.NET Core conventions over micro-optimizations

C# 14 Features Evaluated but Not Used:

- **Extension Members:** Only replaced `Sum()` which is already concise (overengineering)
- **Span optimizations:** No performance bottlenecks in string operations or byte encoding
- **field keyword for validation:** Data Annotations provide better separation of concerns

11. Technology Stack Summary

Layer	Technology
Framework	.NET 10 / ASP.NET Core
Language	C# 14
Database	PostgreSQL 16
ORM	Entity Framework Core 10.0
Message Broker	RabbitMQ 3.13 with Management Plugin
Containerization	Docker & Docker Compose
API Documentation	Scalar.AspNetCore (OpenAPI)
Logging	Serilog with Console and File Sinks
Design Patterns	Clean Architecture, Repository Pattern, Result Pattern

12. Future Enhancements

1. **Caching Layer:** Redis for frequently accessed data
2. **Dead Letter Queue:** Error handling for failed RabbitMQ messages
3. **GraphQL Support:** Alternative to REST API
4. **Real-time Updates:** SignalR for live position updates
5. **Time Series Optimization:** TimescaleDB for historical data
6. **Authentication:** Identity Server or Azure AD B2C
7. **Monitoring:** Application Insights, Prometheus + Grafana
8. **Event Consumers:** Separate microservices subscribing to RabbitMQ events